

Question 1: What is the output of \dt?

```
dvdrental=# \dt
List of relations
Schema | Name | Type | Owner
-----+-----+-----+-----
public | actor | table | postgres
public | address | table | postgres
public | category | table | postgres
public | city | table | postgres
public | country | table | postgres
public | customer | table | postgres
public | film | table | postgres
public | film_actor | table | postgres
public | film_category | table | postgres
public | inventory | table | postgres
public | language | table | postgres
public | payment | table | postgres
public | payment_p2007_01 | table | postgres
public | payment_p2007_02 | table | postgres
public | payment_p2007_03 | table | postgres
public | payment_p2007_04 | table | postgres
public | payment_p2007_05 | table | postgres
public | payment_p2007_06 | table | postgres
public | rental | table | postgres
public | staff | table | postgres
public | store | table | postgres
(21 rows)
```

Question 2: What is the schema for the customer table?

```
dvdrental=# \d customer
Table "public.customer"
Column | Type | Modifiers
-----+-----+-----
customer_id | integer | not null default nextval('customer_customer_id_seq'::regclass)
store_id | smallint | not null
first_name | character varying(45) | not null
last_name | character varying(45) | not null
email | character varying(50) | 
address_id | smallint | not null
activebool | boolean | not null default true
create_date | date | not null default ('now'::text)::date
last_update | timestamp without time zone | default now()
active | integer | 
Indexes:
    "customer_pkey" PRIMARY KEY, btree (customer_id)
    "idx_fk_address_id" btree (address_id)
    "idx_fk_store_id" btree (store_id)
    "idx_last_name" btree (last_name)
Foreign-key constraints:
    "customer_address_id_fkey" FOREIGN KEY (address_id) REFERENCES address(address_id) ON UPDATE CASCADE ON DELETE RESTRICT
    "customer_store_id_fkey" FOREIGN KEY (store_id) REFERENCES store(store_id) ON UPDATE CASCADE ON DELETE RESTRICT
Referenced by:
    TABLE "payment" CONSTRAINT "payment_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id) ON UPDATE CASCADE ON DELETE RESTRICT
    TABLE "payment_p2007_01" CONSTRAINT "payment_p2007_01_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id)
    TABLE "payment_p2007_02" CONSTRAINT "payment_p2007_02_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id)
    TABLE "payment_p2007_03" CONSTRAINT "payment_p2007_03_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id)
    TABLE "payment_p2007_04" CONSTRAINT "payment_p2007_04_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id)
    TABLE "payment_p2007_05" CONSTRAINT "payment_p2007_05_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id)
    TABLE "payment_p2007_06" CONSTRAINT "payment_p2007_06_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id)
    TABLE "rental" CONSTRAINT "rental_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(customer_id) ON UPDATE CASCADE ON DELETE RESTRICT
Triggers:
    last_updated BEFORE UPDATE ON customer FOR EACH ROW EXECUTE PROCEDURE last_updated()
```

Question 3: What similarities do you see in the explain plans for these 3 queries?

All of them use "seq scan" on a chosen table for each case. Two of them use "Append" and "Filter" to filter and append data that match conditions.

```
dvdrental=# EXPLAIN SELECT customer_id, first_name, last_name FROM customer;
               QUERY PLAN
-----
Seq Scan on customer  (cost=0.00..14.99 rows=599 width=17)
(1 row)
```

```
dvdrental=# EXPLAIN SELECT customer_id,
dvdrental-#         amount,
dvdrental-#         payment_date
dvdrental-# FROM payment
dvdrental-# WHERE amount <= 1
dvdrental-#        OR amount >= 8;
               QUERY PLAN
-----
Result  (cost=0.00..420.63 rows=5178 width=19)
-> Append  (cost=0.00..420.63 rows=5178 width=19)
-> Seq Scan on payment  (cost=0.00..29.95 rows=739 width=21)
    Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
-> Seq Scan on payment_p2007_01 payment  (cost=0.00..26.36 rows=266 width=18)
    Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
-> Seq Scan on payment_p2007_02 payment  (cost=0.00..51.68 rows=531 width=18)
    Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
-> Seq Scan on payment_p2007_03 payment  (cost=0.00..126.66 rows=1268 width=18)
    Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
-> Seq Scan on payment_p2007_04 payment  (cost=0.00..151.31 rows=1557 width=18)
    Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
-> Seq Scan on payment_p2007_05 payment  (cost=0.00..4.73 rows=78 width=17)
    Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
-> Seq Scan on payment_p2007_06 payment  (cost=0.00..29.95 rows=739 width=21)
    Filter: ((amount <= 1::numeric) OR (amount >= 8::numeric))
(16 rows)
```

```
dvdrental=# explain SELECT customer_id,
dvdrental-#         payment_id,
dvdrental-#         amount
dvdrental-# FROM payment
dvdrental-# WHERE amount BETWEEN 5 AND 9;
               QUERY PLAN
-----
Result  (cost=0.00..420.63 rows=3600 width=14)
-> Append  (cost=0.00..420.63 rows=3600 width=14)
-> Seq Scan on payment  (cost=0.00..29.95 rows=7 width=17)
    Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
-> Seq Scan on payment_p2007_01 payment  (cost=0.00..26.36 rows=242 width=14)
    Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
-> Seq Scan on payment_p2007_02 payment  (cost=0.00..51.68 rows=506 width=14)
    Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
-> Seq Scan on payment_p2007_03 payment  (cost=0.00..126.66 rows=1290 width=14)
    Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
-> Seq Scan on payment_p2007_04 payment  (cost=0.00..151.31 rows=1535 width=14)
    Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
-> Seq Scan on payment_p2007_05 payment  (cost=0.00..4.73 rows=13 width=13)
    Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
-> Seq Scan on payment_p2007_06 payment  (cost=0.00..29.95 rows=7 width=17)
    Filter: ((amount >= 5::numeric) AND (amount <= 9::numeric))
(16 rows)
```

Question 4: What is the difference between the plans for the Partitioned table and the union query? Why do you think this difference exists?

The biggest difference is that union query plan includes HashAggregate two times while partitioned table query includes HashAggregate only one time. This is probably because for union query applies HashAggregate to each table and then applies HashAggregate to combined table. This makes union query much more expensive than partitioned table.

Union of 2 tables

```
dvdrental=# explain SELECT u.customer_id,
dvdrental=#           sum(u.amount)
dvdrental=# FROM
dvdrental=#   ( SELECT *
dvdrental=#     FROM payment_p2007_01
dvdrental=#     UNION SELECT *
dvdrental=#     FROM payment_p2007_02) u
dvdrental=# WHERE u.payment_date <= '2007-02-01 00:00:00'::TIMESTAMP WITHOUT TIME
dvdrental=# GROUP BY u.customer_id ;

               QUERY PLAN
-----
HashAggregate  (cost=127.26..129.76 rows=200 width=13)
  -> HashAggregate (cost=98.31..109.89 rows=1158 width=28)
    -> Append (cost=0.00..80.94 rows=1158 width=28)
      -> Seq Scan on payment_p2007_01 (cost=0.00..23.46 rows=1157 width=28)
            Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
      -> Seq Scan on payment_p2007_02 (cost=0.00..45.90 rows=1 width=28)
            Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
(7 rows)
```

Partitioned Table

```
dvdrental=# explain SELECT customer_id,
dvdrental=#           sum(amount)
dvdrental=# FROM payment
dvdrental=# WHERE payment_date <= '2007-02-01 00:00:00'::TIMESTAMP WITHOUT TIME
dvdrental=# GROUP BY customer_id ;

               QUERY PLAN
-----
HashAggregate  (cost=103.99..106.49 rows=200 width=11)
  -> Append (cost=0.00..95.99 rows=1601 width=11)
    -> Seq Scan on payment (cost=0.00..26.62 rows=443 width=13)
          Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
    -> Seq Scan on payment_p2007_01 payment (cost=0.00..23.46 rows=1157 width=10)
          Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
    -> Seq Scan on payment_p2007_02 payment (cost=0.00..45.90 rows=1 width=10)
          Filter: (payment_date <= '2007-02-01 00:00:00'::timestamp without time zone)
(8 rows)
```

Question 5: What join algorithm is used for the inner join?

Hash Join

```

dvdrental=# explain SELECT customer.customer_id,
dvdrental=#         first_name,
dvdrental=#         last_name,
dvdrental=#         email,
dvdrental=#         amount,
dvdrental=#         payment_date
dvdrental=# FROM customer
dvdrental=# INNER JOIN payment ON payment.customer_id = customer.customer_id;
                                QUERY PLAN
-----
Hash Join  (cost=22.48..606.82 rows=18709 width=65)
  Hash Cond: (public.payment.customer_id = customer.customer_id)
    -> Append  (cost=0.00..327.09 rows=18709 width=18)
      -> Seq Scan on payment  (cost=0.00..23.30 rows=1330 width=21)
      -> Seq Scan on payment_p2007_01 payment  (cost=0.00..20.57 rows=1157 width=18)
      -> Seq Scan on payment_p2007_02 payment  (cost=0.00..40.12 rows=2312 width=18)
      -> Seq Scan on payment_p2007_03 payment  (cost=0.00..98.44 rows=5644 width=18)
      -> Seq Scan on payment_p2007_04 payment  (cost=0.00..117.54 rows=6754 width=18)
      -> Seq Scan on payment_p2007_05 payment  (cost=0.00..3.82 rows=182 width=17)
      -> Seq Scan on payment_p2007_06 payment  (cost=0.00..23.30 rows=1330 width=21)
    -> Hash  (cost=14.99..14.99 rows=599 width=49)
      -> Seq Scan on customer  (cost=0.00..14.99 rows=599 width=49)
(12 rows)

```